

# Loading CIFAR-10

```
In [ ]: import tensorflow as tf
        from tensorflow.keras import layers, models, datasets
        from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
        from tensorflow.keras.regularizers import l2
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt

In [ ]: # Load and preprocess the CIFAR-10 dataset
        (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data(
        train_images = train_images / 255.0
        test_images = test_images / 255.0
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 4s 0us/step

## Model Building

```
In [ ]: # Define models
        MLP_BN_model = models.Sequential([
            layers.Flatten(input_shape=(32, 32, 3)),
            layers.Dense(512),
            layers.BatchNormalization(),
            layers.Activation('relu'),
            layers.Dropout(0.2),
            layers.Dense(256),
            layers.BatchNormalization(),
            layers.Activation('relu'),
            layers.Dropout(0.2),
            layers.Dense(10, activation='softmax')
        ])

In [ ]: MLP_L2_model = models.Sequential([
            layers.Flatten(input_shape=(32, 32, 3)),
            layers.Dense(512, kernel_regularizer=l2(0.001), activation='relu'),
            layers.Dropout(0.2),
            layers.Dense(10, activation='softmax')
        ])

In [ ]: Simple_CNN_model = models.Sequential([
            layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, (3, 3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Flatten(),
            layers.Dense(64, activation='relu'),
            layers.Dense(10, activation='softmax')
        ])

In [ ]: CNN_DA_Dropout_model = models.Sequential([
            layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
```

```

layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dropout(0.5),
layers.Dense(10, activation='softmax')
])

```

```

In [ ]: ResNet_model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.BatchNormalization(),
    layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(10, activation='softmax')
])

```

## Model Training

```

In [ ]: # Compile all models
models = [MLP_BN_model, MLP_L2_model, Simple_CNN_model, CNN_DA_Dropout_model, ResNet_model]
model_names = ["MLP with BatchNorm", "MLP with L2", "Simple CNN", "CNN with DA and Dropout", "ResNet"]

i=0
for model in models:
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
    print(model_names[i]+" Model Architecture : \n")
    print(model.summary())
    i+=1

# Train all models and store results
results = []

for idx, (model, name) in enumerate(zip(models, model_names), start=1):
    print(f"Training Model {idx}: {name}...")
    history = model.fit(train_images, train_labels, epochs=10, validation_data=(test_images, test_labels))

    # Evaluate the model
    test_loss, test_acc = model.evaluate(test_images, test_labels)
    print(f'Test accuracy for {name}:', test_acc)
    results.append((history.history, test_loss, test_acc))

```

MLP with BatchNorm Model Architecture :

Model: "sequential"

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| flatten (Flatten)                           | (None, 3072) | 0       |
| dense (Dense)                               | (None, 512)  | 1573376 |
| batch_normalization (Batch Normalization)   | (None, 512)  | 2048    |
| activation (Activation)                     | (None, 512)  | 0       |
| dropout (Dropout)                           | (None, 512)  | 0       |
| dense_1 (Dense)                             | (None, 256)  | 131328  |
| batch_normalization_1 (Batch Normalization) | (None, 256)  | 1024    |
| activation_1 (Activation)                   | (None, 256)  | 0       |
| dropout_1 (Dropout)                         | (None, 256)  | 0       |
| dense_2 (Dense)                             | (None, 10)   | 2570    |

Total params: 1710346 (6.52 MB)  
Trainable params: 1708810 (6.52 MB)  
Non-trainable params: 1536 (6.00 KB)

None

MLP with L2 Model Architecture :

Model: "sequential\_1"

| Layer (type)        | Output Shape | Param # |
|---------------------|--------------|---------|
| flatten_1 (Flatten) | (None, 3072) | 0       |
| dense_3 (Dense)     | (None, 512)  | 1573376 |
| dropout_2 (Dropout) | (None, 512)  | 0       |
| dense_4 (Dense)     | (None, 10)   | 5130    |

Total params: 1578506 (6.02 MB)  
Trainable params: 1578506 (6.02 MB)  
Non-trainable params: 0 (0.00 Byte)

None

Simple CNN Model Architecture :

Model: "sequential\_2"

| Layer (type) | Output Shape | Param # |
|--------------|--------------|---------|
|--------------|--------------|---------|

|                                |                    |        |
|--------------------------------|--------------------|--------|
| conv2d (Conv2D)                | (None, 30, 30, 32) | 896    |
| max_pooling2d (MaxPooling2D)   | (None, 15, 15, 32) | 0      |
| conv2d_1 (Conv2D)              | (None, 13, 13, 64) | 18496  |
| max_pooling2d_1 (MaxPooling2D) | (None, 6, 6, 64)   | 0      |
| flatten_2 (Flatten)            | (None, 2304)       | 0      |
| dense_5 (Dense)                | (None, 64)         | 147520 |
| dense_6 (Dense)                | (None, 10)         | 650    |

=====  
Total params: 167562 (654.54 KB)  
Trainable params: 167562 (654.54 KB)  
Non-trainable params: 0 (0.00 Byte)

---

None

CNN with DA and Dropout Model Architecture :

Model: "sequential\_3"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_2 (Conv2D)              | (None, 30, 30, 32) | 896     |
| max_pooling2d_2 (MaxPooling2D) | (None, 15, 15, 32) | 0       |
| conv2d_3 (Conv2D)              | (None, 13, 13, 64) | 18496   |
| max_pooling2d_3 (MaxPooling2D) | (None, 6, 6, 64)   | 0       |
| flatten_3 (Flatten)            | (None, 2304)       | 0       |
| dense_7 (Dense)                | (None, 128)        | 295040  |
| dropout_3 (Dropout)            | (None, 128)        | 0       |
| dense_8 (Dense)                | (None, 10)         | 1290    |

=====  
Total params: 315722 (1.20 MB)  
Trainable params: 315722 (1.20 MB)  
Non-trainable params: 0 (0.00 Byte)

---

None

ResNet Model Architecture :

Model: "sequential\_4"

| Layer (type)      | Output Shape       | Param # |
|-------------------|--------------------|---------|
| conv2d_4 (Conv2D) | (None, 30, 30, 32) | 896     |

|   |                    |        |
|---|--------------------|--------|
| batch_normalization_2 (Batch Normalization) | (None, 30, 30, 32) | 128    |
| conv2d_5 (Conv2D)                           | (None, 30, 30, 32) | 9248   |
| batch_normalization_3 (Batch Normalization) | (None, 30, 30, 32) | 128    |
| max_pooling2d_4 (MaxPooling2D)              | (None, 15, 15, 32) | 0      |
| conv2d_6 (Conv2D)                           | (None, 15, 15, 64) | 18496  |
| batch_normalization_4 (Batch Normalization) | (None, 15, 15, 64) | 256    |
| conv2d_7 (Conv2D)                           | (None, 15, 15, 64) | 36928  |
| batch_normalization_5 (Batch Normalization) | (None, 15, 15, 64) | 256    |
| max_pooling2d_5 (MaxPooling2D)              | (None, 7, 7, 64)   | 0      |
| conv2d_8 (Conv2D)                           | (None, 7, 7, 128)  | 73856  |
| batch_normalization_6 (Batch Normalization) | (None, 7, 7, 128)  | 512    |
| conv2d_9 (Conv2D)                           | (None, 7, 7, 128)  | 147584 |
| batch_normalization_7 (Batch Normalization) | (None, 7, 7, 128)  | 512    |
| max_pooling2d_6 (MaxPooling2D)              | (None, 3, 3, 128)  | 0      |
| flatten_4 (Flatten)                         | (None, 1152)       | 0      |
| dense_9 (Dense)                             | (None, 128)        | 147584 |
| dropout_4 (Dropout)                         | (None, 128)        | 0      |
| dense_10 (Dense)                            | (None, 10)         | 1290   |

```

=====
Total params: 437674 (1.67 MB)
Trainable params: 436778 (1.67 MB)
Non-trainable params: 896 (3.50 KB)

```

---

None

Training Model 1: MLP with BatchNorm...

Epoch 1/10

1563/1563 [=====] - 13s 6ms/step - loss: 1.7347 - accuracy: 0.3832 - val\_loss: 2.1473 - val\_accuracy: 0.3034

Epoch 2/10

1563/1563 [=====] - 8s 5ms/step - loss: 1.5232 - accuracy: 0.4553 - val\_loss: 1.5415 - val\_accuracy: 0.4547

Epoch 3/10

1563/1563 [=====] - 8s 5ms/step - loss: 1.4411 - accuracy:

0.4864 - val\_loss: 1.6450 - val\_accuracy: 0.4087  
Epoch 4/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.3848 - accuracy:  
0.5071 - val\_loss: 1.5184 - val\_accuracy: 0.4530  
Epoch 5/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.3353 - accuracy:  
0.5241 - val\_loss: 1.5725 - val\_accuracy: 0.4592  
Epoch 6/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.2948 - accuracy:  
0.5386 - val\_loss: 1.4949 - val\_accuracy: 0.4582  
Epoch 7/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.2556 - accuracy:  
0.5538 - val\_loss: 1.3717 - val\_accuracy: 0.5053  
Epoch 8/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.2233 - accuracy:  
0.5654 - val\_loss: 1.4717 - val\_accuracy: 0.4853  
Epoch 9/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.1933 - accuracy:  
0.5770 - val\_loss: 1.3859 - val\_accuracy: 0.5016  
Epoch 10/10  
1563/1563 [=====] - 8s 5ms/step - loss: 1.1593 - accuracy:  
0.5863 - val\_loss: 1.3731 - val\_accuracy: 0.5116  
313/313 [=====] - 1s 3ms/step - loss: 1.3731 - accuracy: 0.5  
116  
Test accuracy for MLP with BatchNorm: 0.5116000175476074  
Training Model 2: MLP with L2...  
Epoch 1/10  
1563/1563 [=====] - 6s 4ms/step - loss: 2.1863 - accuracy:  
0.2863 - val\_loss: 1.8867 - val\_accuracy: 0.3415  
Epoch 2/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.9206 - accuracy:  
0.3199 - val\_loss: 1.8613 - val\_accuracy: 0.3676  
Epoch 3/10  
1563/1563 [=====] - 5s 3ms/step - loss: 1.9101 - accuracy:  
0.3265 - val\_loss: 1.8593 - val\_accuracy: 0.3464  
Epoch 4/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.8926 - accuracy:  
0.3337 - val\_loss: 1.8471 - val\_accuracy: 0.3685  
Epoch 5/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.8924 - accuracy:  
0.3349 - val\_loss: 1.7962 - val\_accuracy: 0.3881  
Epoch 6/10  
1563/1563 [=====] - 5s 3ms/step - loss: 1.8865 - accuracy:  
0.3353 - val\_loss: 1.8175 - val\_accuracy: 0.3823  
Epoch 7/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.8876 - accuracy:  
0.3375 - val\_loss: 1.7869 - val\_accuracy: 0.3835  
Epoch 8/10  
1563/1563 [=====] - 5s 3ms/step - loss: 1.8756 - accuracy:  
0.3406 - val\_loss: 1.8704 - val\_accuracy: 0.3443  
Epoch 9/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.8839 - accuracy:  
0.3344 - val\_loss: 1.7968 - val\_accuracy: 0.3650  
Epoch 10/10  
1563/1563 [=====] - 5s 3ms/step - loss: 1.8740 - accuracy:  
0.3367 - val\_loss: 1.7492 - val\_accuracy: 0.3970  
313/313 [=====] - 1s 2ms/step - loss: 1.7492 - accuracy: 0.3  
970  
Test accuracy for MLP with L2: 0.3970000147819519  
Training Model 3: Simple CNN...

Epoch 1/10  
1563/1563 [=====] - 11s 4ms/step - loss: 1.4572 - accuracy: 0.4779 - val\_loss: 1.2895 - val\_accuracy: 0.5335

Epoch 2/10  
1563/1563 [=====] - 7s 5ms/step - loss: 1.1258 - accuracy: 0.6071 - val\_loss: 1.0722 - val\_accuracy: 0.6217

Epoch 3/10  
1563/1563 [=====] - 6s 4ms/step - loss: 0.9933 - accuracy: 0.6532 - val\_loss: 1.0358 - val\_accuracy: 0.6477

Epoch 4/10  
1563/1563 [=====] - 7s 5ms/step - loss: 0.9107 - accuracy: 0.6848 - val\_loss: 0.9552 - val\_accuracy: 0.6718

Epoch 5/10  
1563/1563 [=====] - 6s 4ms/step - loss: 0.8413 - accuracy: 0.7080 - val\_loss: 0.9666 - val\_accuracy: 0.6644

Epoch 6/10  
1563/1563 [=====] - 7s 5ms/step - loss: 0.7812 - accuracy: 0.7292 - val\_loss: 0.9419 - val\_accuracy: 0.6815

Epoch 7/10  
1563/1563 [=====] - 6s 4ms/step - loss: 0.7333 - accuracy: 0.7445 - val\_loss: 0.9341 - val\_accuracy: 0.6884

Epoch 8/10  
1563/1563 [=====] - 7s 5ms/step - loss: 0.6869 - accuracy: 0.7604 - val\_loss: 0.9366 - val\_accuracy: 0.6888

Epoch 9/10  
1563/1563 [=====] - 6s 4ms/step - loss: 0.6418 - accuracy: 0.7754 - val\_loss: 0.9704 - val\_accuracy: 0.6838

Epoch 10/10  
1563/1563 [=====] - 7s 5ms/step - loss: 0.6015 - accuracy: 0.7897 - val\_loss: 0.9458 - val\_accuracy: 0.6909  
313/313 [=====] - 1s 2ms/step - loss: 0.9458 - accuracy: 0.6909

Test accuracy for Simple CNN: 0.6909000277519226

Training Model 4: CNN with DA and Dropout...

Epoch 1/10  
1563/1563 [=====] - 9s 5ms/step - loss: 1.6409 - accuracy: 0.3993 - val\_loss: 1.2958 - val\_accuracy: 0.5445

Epoch 2/10  
1563/1563 [=====] - 6s 4ms/step - loss: 1.3169 - accuracy: 0.5320 - val\_loss: 1.1419 - val\_accuracy: 0.5916

Epoch 3/10  
1563/1563 [=====] - 7s 5ms/step - loss: 1.1903 - accuracy: 0.5821 - val\_loss: 1.0418 - val\_accuracy: 0.6381

Epoch 4/10  
1563/1563 [=====] - 7s 4ms/step - loss: 1.1090 - accuracy: 0.6100 - val\_loss: 0.9715 - val\_accuracy: 0.6582

Epoch 5/10  
1563/1563 [=====] - 7s 5ms/step - loss: 1.0510 - accuracy: 0.6313 - val\_loss: 0.9415 - val\_accuracy: 0.6731

Epoch 6/10  
1563/1563 [=====] - 7s 5ms/step - loss: 1.0090 - accuracy: 0.6480 - val\_loss: 0.9243 - val\_accuracy: 0.6779

Epoch 7/10  
1563/1563 [=====] - 7s 4ms/step - loss: 0.9588 - accuracy: 0.6638 - val\_loss: 0.9108 - val\_accuracy: 0.6806

Epoch 8/10  
1563/1563 [=====] - 8s 5ms/step - loss: 0.9260 - accuracy: 0.6760 - val\_loss: 0.8999 - val\_accuracy: 0.6857

Epoch 9/10  
1563/1563 [=====] - 7s 4ms/step - loss: 0.8957 - accuracy:

```

0.6848 - val_loss: 0.8889 - val_accuracy: 0.6946
Epoch 10/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.8745 - accuracy:
0.6923 - val_loss: 0.8904 - val_accuracy: 0.6956
313/313 [=====] - 1s 2ms/step - loss: 0.8904 - accuracy: 0.6
956
Test accuracy for CNN with DA and Dropout: 0.6955999732017517
Training Model 5: ResNet...
Epoch 1/10
1563/1563 [=====] - 18s 9ms/step - loss: 1.5953 - accuracy:
0.4298 - val_loss: 1.1129 - val_accuracy: 0.6007
Epoch 2/10
1563/1563 [=====] - 12s 8ms/step - loss: 1.1086 - accuracy:
0.6112 - val_loss: 1.0145 - val_accuracy: 0.6411
Epoch 3/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.9156 - accuracy:
0.6877 - val_loss: 0.9534 - val_accuracy: 0.6802
Epoch 4/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.7776 - accuracy:
0.7380 - val_loss: 0.8449 - val_accuracy: 0.7255
Epoch 5/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.6726 - accuracy:
0.7745 - val_loss: 0.7478 - val_accuracy: 0.7575
Epoch 6/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.5857 - accuracy:
0.8028 - val_loss: 0.7808 - val_accuracy: 0.7399
Epoch 7/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.5164 - accuracy:
0.8273 - val_loss: 0.6081 - val_accuracy: 0.8064
Epoch 8/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.4532 - accuracy:
0.8476 - val_loss: 0.6316 - val_accuracy: 0.7951
Epoch 9/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.4013 - accuracy:
0.8648 - val_loss: 0.5905 - val_accuracy: 0.8172
Epoch 10/10
1563/1563 [=====] - 12s 7ms/step - loss: 0.3594 - accuracy:
0.8782 - val_loss: 0.6309 - val_accuracy: 0.8137
313/313 [=====] - 1s 3ms/step - loss: 0.6309 - accuracy: 0.8
137
Test accuracy for ResNet: 0.8137000203132629

```

```

In [21]: import matplotlib.pyplot as plt

# Function to plot Learning curves
def plot_learning_curves(history, title):
    plt.figure(figsize=(12, 6))

    # Plot training & validation accuracy values
    plt.subplot(1, 2, 1)
    plt.plot(history['accuracy'])
    plt.plot(history['val_accuracy'])
    plt.title(title + ' - Model Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend(['Train', 'Validation'], loc='lower right')
    plt.grid(True)

    # Plot training & validation loss values
    plt.subplot(1, 2, 2)

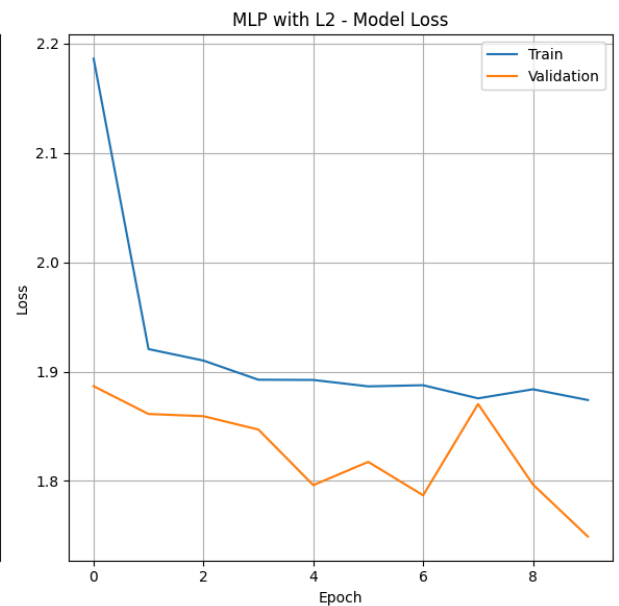

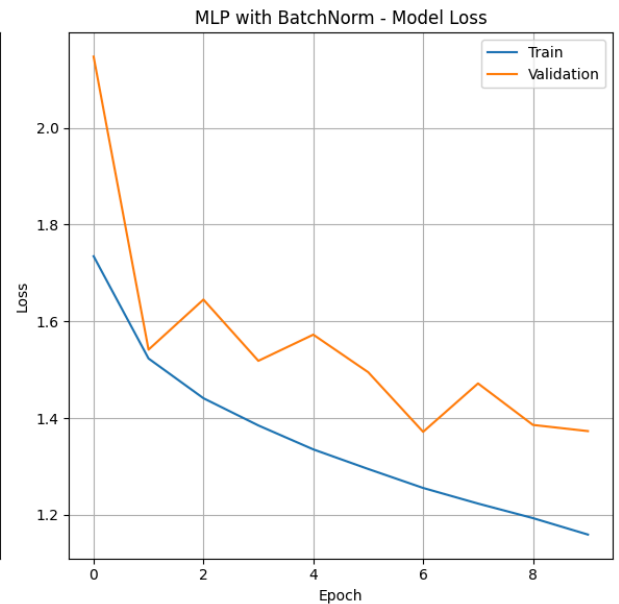
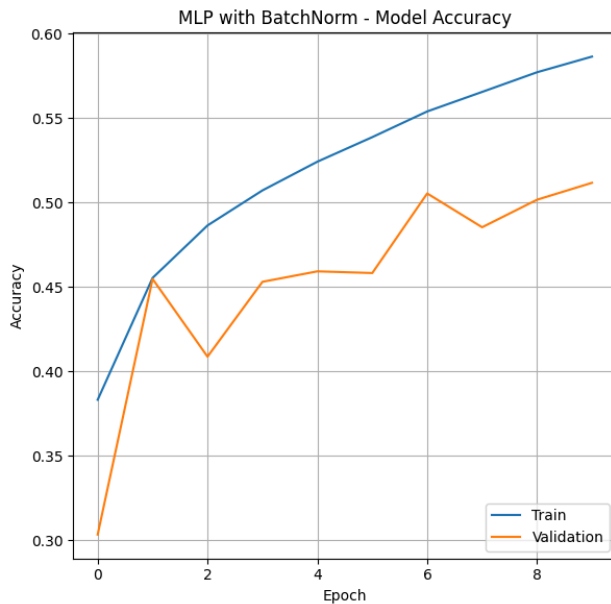
```

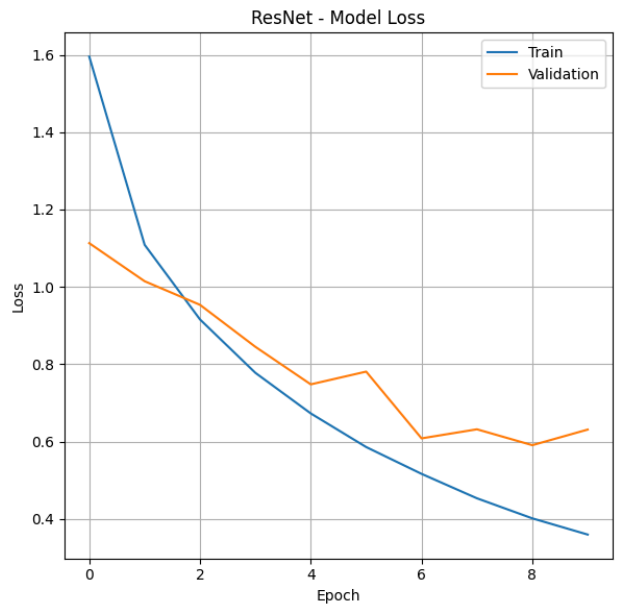
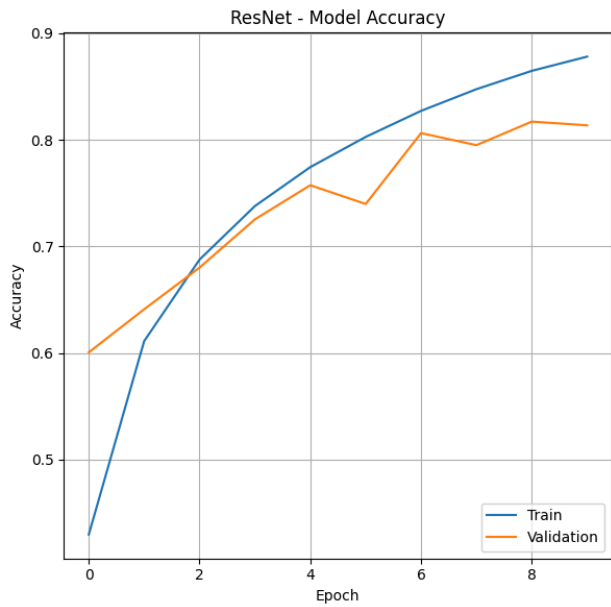
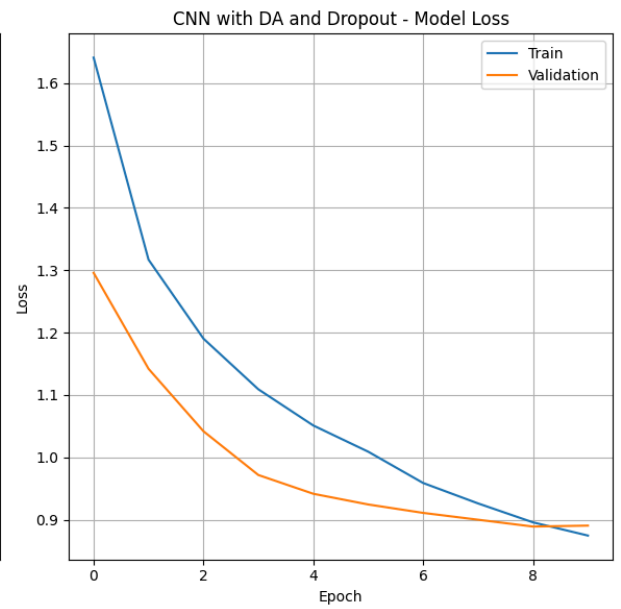

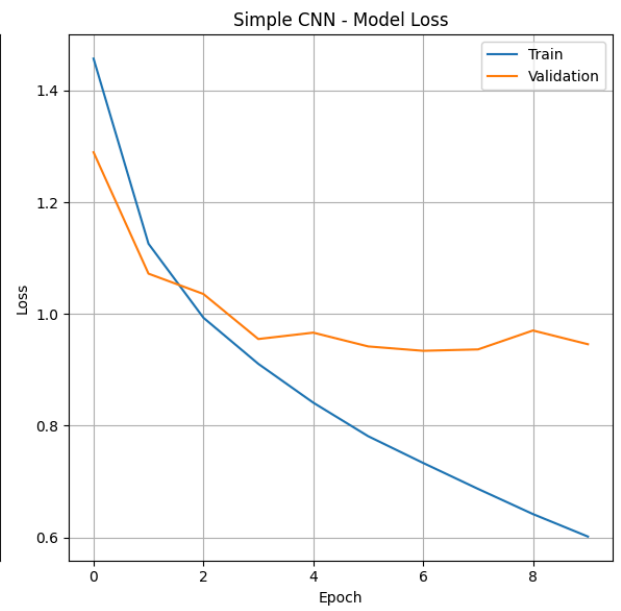



```
plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title(title + ' - Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.grid(True)
```

```
plt.tight_layout()
plt.show()
```

```
# Plot learning curves for each model
for history, name in zip(results, model_names):
    plot_learning_curves(history[0], name)
```





# Model Comparison

```
In [ ]: # Calculate metrics for each model
accuracy_scores = []
precision_scores = []
recall_scores = []
f1_scores = []

for name, model in zip(model_names, models):
    pred = np.argmax(model.predict(test_images), axis=1)
    accuracy = accuracy_score(test_labels, pred)
    precision = precision_score(test_labels, pred, average='macro')
    recall = recall_score(test_labels, pred, average='macro')
    f1 = f1_score(test_labels, pred, average='macro')

    accuracy_scores.append(accuracy)
    precision_scores.append(precision)
    recall_scores.append(recall)
    f1_scores.append(f1)

# Create a DataFrame to store the metrics
metrics_df = pd.DataFrame({
    'Model': model_names,
    'Accuracy': accuracy_scores,
    'Precision': precision_scores,
    'Recall': recall_scores,
    'F1 Score': f1_scores
})

# Print the metrics DataFrame
print("Metrics for each model:")
display(metrics_df)
```

```
313/313 [=====] - 1s 2ms/step
313/313 [=====] - 1s 2ms/step
313/313 [=====] - 1s 2ms/step
313/313 [=====] - 1s 2ms/step
313/313 [=====] - 1s 2ms/step
Metrics for each model:
```

|   | Model                   | Accuracy | Precision | Recall | F1 Score |
|---|-------------------------|----------|-----------|--------|----------|
| 0 | MLP with BatchNorm      | 0.5116   | 0.536297  | 0.5116 | 0.511579 |
| 1 | MLP with L2             | 0.3970   | 0.390330  | 0.3970 | 0.379442 |
| 2 | Simple CNN              | 0.6909   | 0.692682  | 0.6909 | 0.687796 |
| 3 | CNN with DA and Dropout | 0.6956   | 0.695649  | 0.6956 | 0.692129 |
| 4 | ResNet                  | 0.8137   | 0.819477  | 0.8137 | 0.814185 |

```
In [ ]: import matplotlib.pyplot as plt

# Plotting the metrics
metrics_df.set_index('Model', inplace=True)

plt.figure(figsize=(12, 8))
```

```

metrics_df.plot(kind='bar', colormap='viridis', alpha=0.8)
plt.title('Performance Metrics Comparison')
plt.xlabel('Model')
plt.ylabel('Score')
plt.xticks(rotation=45, ha='right')
plt.legend(loc='upper center', bbox_to_anchor=(0.5, -0.25), ncol=len(metrics_df.columns))

plt.tight_layout()
plt.show()

```

<Figure size 1200x800 with 0 Axes>

